

Assignment Project Exam Help

Operating Systems and Concurrency

Lecture 18: File Systems I - Disk Scheduling

G52OSC

<https://powcoder.com>

Add WeChat powcoder
Geert De Maere and Isaac Triguero
{Geert.DeMaere,Isaac.Triguero}@Nottingham.ac.uk

University Of Nottingham
United Kingdom

2018

Assignment Project Exam Help

- Construction of hard disks
- Accessing hard disks
- Disk scheduling

<https://powcoder.com>

Add WeChat powcoder

¹Slides partially based on slides by Colin Higgins and Jon Garibaldi

Hard Disks

Construction of Hard Disks

Assignment Project Exam Help

- Disks are constructed as multiple aluminium/glass **platters** covered with **magnetisable material**
 - **Read/write heads** fly just above the surface (0.2 – 0.07mm) and are connected to a single **disk arm** controlled by a single **actuator**
 - Data is stored on both sides
 - Common **diameters** range from 1.8 to 3.5 inches
 - Hard disks **rotate** at a **constant speed** (i.e., speed on the inside less than on the outside \leftrightarrow CD-ROMS)
- A **disk controller** sits between the CPU and the drive
- Hard disks are currently about **4 orders of magnitude slower than main memory** \Rightarrow how can we reduce the impact of this?

<https://powcoder.com>

Add WeChat powcoder

Hard Disks

Construction of Hard Disks

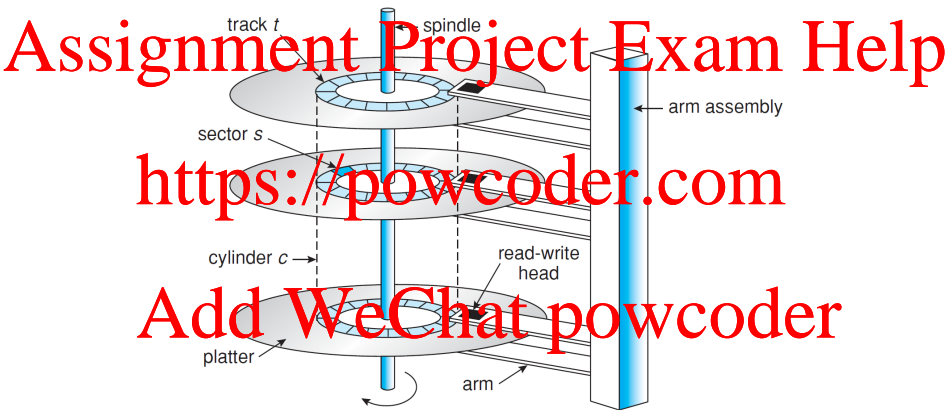


Figure: Construction of a Hard Disk (Silberschatz)

Hard Disks

Low Level Format

Assignment Project Exam Help

- Disks are organised into:
 - **Cylinders**: a collection of tracks in the same relative position to the spindle
 - **Tracks**: a concentric circle on a single platter side
 - **Sectors**: segments of a track (usually 512B or 4KB in size)
- Sectors usually have an equal number of bytes in them, consisting of a **preamble**, **data**, and an **error correcting code**
- The **number of sectors increases** from the inner side of the disk to the outside

<https://powcoder.com>
Add WeChat powcoder



Figure: Disk Sector

Hard Disks

Organisation of Hard Disks

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

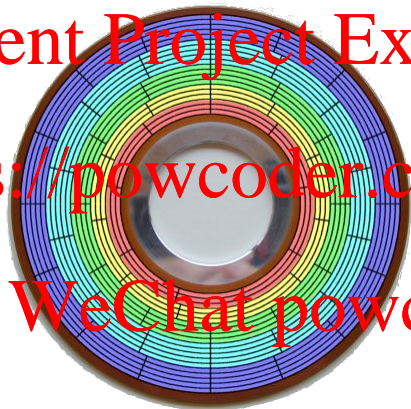


Figure: Disk Layout: Cylinder skew (Source: www.pcguide.com/)

Hard Disks

Organisation of Hard Disks

Assignment Project Exam Help

- Disks usually have a **cylinder skew**: i.e., an **offset** is added to sector 0 in adjacent tracks to account for the seek time
- In the past, consecutive **disk sectors were interleaved** to account for transfer time
- Note that as a result of this low-level formatting, disk capacity is reduced (size of preamble, ECC, etc.)

<https://powcoder.com>

Add WeChat powcoder

Hard Disks

Access Times

- **Access time** = seek time + rotational delay + transfer time
 - **Seek time** = time needed to move the arm to the cylinder (dominant)
 - **Rotational latency** = time before the sector appears under the head (on average half the rotation time)
 - **Transfer time** = time to transfer the data




Figure: Access time to Disk (Source: www.studiodaily.com/)

Hard Disks

Access Times

- Multiple requests may be happening at the same time (concurrently). Thus, access time may be increased by a **queuing time**
- In this scenario, dominance of seek time leaves room for **optimisation** by carefully considering the order of read operations

<https://powcoder.com>

Disk Delay	Queuing	Seek Time	Rotational Latency	Transfer Time
------------	---------	-----------	--------------------	---------------

Add WeChat powcoder



Hard Disks

Access Times

Assignment Project Exam Help

- The **estimated seek time** (i.e., to move the arm from one track to another) is approximated by:

$$T_s = n \times m + s \quad (1)$$

- In which T_s denotes the estimated seek time, n the **number of tracks** to be crossed, m the **crossing time per track**, and s any **additional startup delay**

Hard Disks

Access Times

Assignment Project Exam Help

- Let us assume a disk that rotates at 3600 rpm (common rotation speeds are between 3600 and 15000 rpm)

- One rotation takes approx. 16.7ms ($\Rightarrow \frac{1}{x} = \frac{3600}{60 \times 1000} \Leftrightarrow x = \frac{60000}{3600}$)
- The average **rotational latency** (T_r) is then $\frac{16.7}{2} \approx 8.3ms$

- Let b denote the **number of bytes transferred**, N the **number of bytes per track**, and rpm the **rotation speed** in rotations per minute, the **transfer time**, T_t , is then given by:

- N bytes take 1 revolution $\Rightarrow \frac{60000}{3600} ms = \frac{ms \text{ per minute}}{rpm}$
- b contiguous bytes takes $\frac{b}{N}$ revolutions

Add WeChat powcoder

$$T_t = \frac{b}{N} \times \frac{ms \text{ per minute}}{rpm} \quad (2)$$

Hard Disks

Access Times: Example

Assignment Project Exam Help

- Read a file of size 256 sectors with:
 - $T_s = 20$ ms (average seek time)
 - 32 sectors/track
- Suppose the file is stored as compact as possible - contiguous, i.e., all sectors on 8 consecutive tracks of 32 sectors each (sequential storage)
 - The first track takes: seek + rotational delay + transfer time
 $\rightarrow 20 + 8.3 + 16.7 = 45ms$
 - Assuming no cylinder skew, and neglecting small seeks between tracks -
 We only need to account for rotational delay + transfer time:
 $8.3 + 16.7 = 25ms$
- The total time is then $45 + 7 \times 25 = 220ms = 0.22$ s

Hard Disks

Access Times: Example

Assignment Project Exam Help

- In case the access is not sequential but at **random for the sectors**, we get:

- $$T_{\text{time per sector}} = T_s + T_r + T_t = 20 + 3.3 + 0.5 = 23.8 \text{ ms}$$

$$T_t = 16.7 \times \frac{1}{32} = 0.5$$

- $$\text{Total time 256 sectors} = 256 \times 28.8 \text{ ms} = 7.37 \text{ s}$$

- It is important to **position the sectors carefully** and avoid **disk fragmentation**.

Add WeChat powcoder

Disk Scheduling

Concepts

Assignment Project Exam Help

- The OS must use the hardware efficiently:
 - The file system can **position/organise files strategically**
 - Having multiple disk requests in a queue allows us to **minimise the arm movement**
- Note that every I/O operation goes through a system call, allowing the **operating system to intercept the request and resequence it**
- If the drive (or the controller) is **free**, the request can be serviced immediately, if not, the request will be **queued**

<https://powcoder.com>

Add WeChat powcoder

Disk Scheduling

Concepts

Assignment Project Exam Help

- In a dynamic situation, several I/O requests will be **made over time** that are kept in a **table of requested sectors per cylinder**
- **Disk scheduling algorithms** determine the order in which disk events are processed
- **None** of the algorithms discussed here are **optimal algorithms**
- Assume a disk with 36 cylinders, numbered 1 to 36

Disk Scheduling

First-Come, First-Served

- **First come first served:** process the requests in the order that they arrive

- Consider the following sequence of disk requests (cylinder locations):

11 1 36 16 34 9 12

- In the order of arrival (FCFS) the total length is:

$|11-1| + |1-36| + |36-16| + |16-34| + |34-9| + |9-12| = 111$




Figure: Head movement for FCFS

Disk Scheduling

Shortest Seek Time First

- **Shortest seek time first** selects the request that is closest to the current head position to reduce head movement

- In the order "shortest seek time first, SSTF" (or shortest job first) we can approx. 50% (for 11 1 36 16 34 9 12):

$$|11-12|+|12-9|+|9-16|+|16-1|+|1-34|+|34-36|=61$$

<https://powcoder.com>

Add WeChat powcoder




Figure: Head movement for shortest seek time

Disk Scheduling

Shortest Seek Time First

Assignment Project Exam Help

- Shortest seek time first could result in **starvation**:
 - The arm stays in the middle of the disk in case of heavy load, edge cylinders are poorly served, the strategy is unfair
 - Continuously arriving requests for the same location could **starve** other regions

Add WeChat powcoder

Disk Scheduling

SCAN

Assignment Project Exam Help

- "lift algorithm, **SCAN**" keep moving in the same direction until end is reached (start upwards):
 - It continues in the current direction, **servicing all pending requests** as it passes over them
 - When it gets to the **last cylinder**, it **reverses direction** and services all the pending requests (until it reaches the first cylinder)
- (Dis-)advantages include:
 - The **upper limit** on the "waiting time" is $2 \times$ number of cylinders, i.e. **no starvation occurs**
 - The **middle cylinders are favoured** if the disk is heavily used (max. wait time is N tracks, $2N$ for the cylinders on the edge)

<https://powcoder.com>

Add WeChat powcoder

Disk Scheduling

SCAN

- “Lift algorithm, SCAN” (for 11 1 36 16 34 9 12):

$$|11-12|+|12-16|+|16-34|+|34-36|+|36-9|+|9-1|=60$$




Figure: Head movement for SCAN

Disk Scheduling

C-SCAN

Assignment Project Exam Help

- Once the outer/inner side of the disk is reached, the **requests at the other end of the disk have been waiting longest**
- SCAN can be improved by using a **circular scan** approach \Rightarrow C-SCAN
 - The disk arm moves in one direction servicing requests
 - When it gets to the last cylinder of the disk, it **reverses direction** but it **does not service requests** on the return journey
 - Once it gets back to the first cylinder it reverses direction, and again services requests
 - It is **fairer** and **equalses response times** across a disk
- The C-SCAN algorithm (for 11 1 36 16 34 9 12):
 $|11-12|+|12-16|+|16-34|+|34-36|+|36-1|+|1-9|=68$

Disk Scheduling

Other SCAN variations: LOOK-SCAN, N-step-SCAN

Assignment Project Exam Help

- Look-SCAN moves to the cylinder containing **the first/last request** (as opposed to the first/last cylinder on the disk with SCAN)
- However, seeks are **cylinder by cylinder**, and one cylinder contains multiple tracks.
- It may happen that the arm “sticks” to a cylinder.
- N-step-SCAN only services N requests every single sweep

<https://powcoder.com>

Add WeChat powcoder

Disk Scheduling

Observations

Assignment Project Exam Help

- Look-SCAN and variations seem to be reasonable choices for the algorithms
- Performance of the algorithms is dependent on the requests/load of the disk
 - One request at a time \Rightarrow FCFS will perform equally well as any other algorithm
- **Optimal algorithms** are difficult to achieve if requests arrive over time!

<https://powcoder.com>

Add WeChat powcoder

Disk scheduling in Unix/Linux

Modifying the disk scheduler

Assignment Project Exam Help

- In Linux, we can modify the disk scheduler by modifying the file: `/sys/block/sda/queue/scheduler`
- We have got three policies:
 - `noop`: this is FCFS
 - `deadline`: N-step-SCAN
 - `cfq`: Complete Fairness Queueing from Linux.
- The one between brackets is the current policy.

<https://powcoder.com>

Add WeChat powcoder

```
pszit@severn:~$ cat /sys/block/sda/queue/scheduler  
noop [deadline] cfq
```


Hard Disks

Driver Caching

Assignment Project Exam Help

- For most current drives, **the time required to seek** a new cylinder is **more than the rotational time** (remember **pre-paging** in this context!)
- It makes sense, therefore, to **read more sectors than actually required**
 - **Read** sectors during the **rotational delay** (i.e. that accidentally pass by)
 - **Modern controllers read multiple sectors** when asked for the data from one sector: track-at-a-time caching

Add WeChat powcoder

Disk scheduling SSD drives

Do we have to do any scheduling at all?

Assignment Project Exam Help

- Solid State Drives (SSDs) have no moving parts and store data using electrical circuits.
 - They don't have T_{seek} or rotational delay!
 - FCFS algorithm is useful in general purposes systems
 - SSTF, SCAN, LOOK-SCAN may reduce performance (no heads to move)

<https://powcoder.com>
Add WeChat powcoder

Summary

Take-Home Message²

Assignment Project Exam Help

- Construction and organisation of hard disks
- Access times of hard disks
- Disk scheduling
- Disk caching

<https://powcoder.com>

Add WeChat powcoder

²Tanenbaum Section 5.4.1 (excluding raid), 5.4.2, 5.4.3

Disk Scheduling

Problem (From Tanenbaum)

Assignment Project Exam Help

Disk requests come in to the disk driver for cylinders 10, 22, 20, 2, 40, 6 and 38, in that order.

- A seek takes 6ms per cylinder.
- How much seek time is needed for: FCFS, SSTF and Look-SCAN (initially moving upward)
- In all cases, the arm is initially at cylinder 20.

Submit your answer at:

<https://b.socrative.com/login/student/>

Room name: G52OSC