

Assignment Project Exam Help

Operating Systems and Concurrency

Lecture 21: File Systems IV

G52OSC

<https://powcoder.com>

Add WeChat [powcoder](https://powcoder.com)
{Geert.DelMaere,Isaac.Triguero}@Nottingham.ac.uk

University Of Nottingham
United Kingdom

2018

Recap

Last Lecture: file system implementations

Assignment Project Exam Help

- **Contiguous implementations** are easy, fast, but result in external fragmentation
- **Linked lists** are sequential, and have block sizes $\neq 2^n$ (page sizes are 2^n)
- **FAT** have block sizes $= 2^n$, but the table becomes prohibitively large
- **I-nodes** are only loaded when the file is open, contain attributes and multiple block levels

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

- Hard and soft links with i-nodes
- File system **paradigms** (which use an underlying file system)
 - Log structured file systems
 - Journaling file systems
 - The Unix/Linux virtual file system

<https://powcoder.com>
Add WeChat powcoder

I-nodes: Sharing files between directories

Hard and Soft Links

Assignment Project Exam Help

- There are two approaches to **share a file**, e.g. between directory B and C, where C is the 'real' owner:

- **Hard links:** maintain two (or multiple) references to the same i-node in B and C. (In Unix: `ln /C/file1 /B/file2`)

<https://powcoder.com>

- **Symbolic links:**

- The owner maintains a reference to the i-node in, e.g., directory C
- The "referencer" maintains a small file (that has its own i-node) that contains the location and name of the shared file in directory C (In Unix: `ln -s /C/file1 /B/file2`)

Add WeChat powcoder

- What is the **best approach**? \Rightarrow both have advantages and disadvantages

I-nodes: Sharing files between directories

Hard Links

Assignment Project Exam Help

- Hard links are the **fastest way of linking files!**
- Disadvantages of hard links:
 - Assume that the owner of the file **deletes** it:
 - If the i-node is also deleted, any hard link will, in the best case, **point to an invalid i-node**
 - If the i-node gets **deleted** and “**recycled**” to point to an other file, the hard links will **point to the wrong file!**
 - The only solution is to **delete the file**, and **leave the i-node intact** if the “**reference count**” is larger than 0 (the original owner of the file still gets “charged” for the space)

<https://powcoder.com>

Add WeChat powcoder

I-nodes: Sharing files between directories

Hard and Soft Links

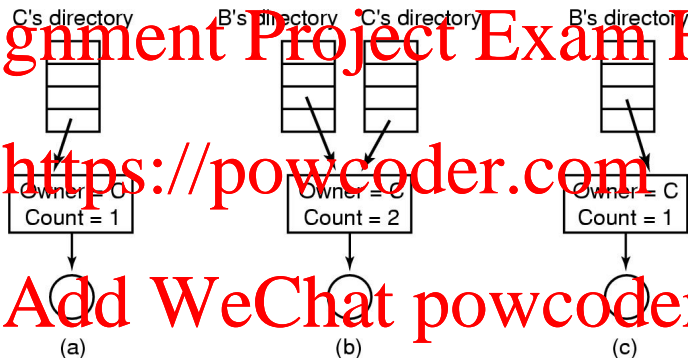


Figure: (a) Before creating a link. (b) After creating a link. (c) after the original owner removes the file (Tanenbaum)

I-nodes

Soft Links

Assignment Project Exam Help

- Disadvantages of soft links:
 - They result in an **extra file lookup** (once the link file has been found, the original file needs to be found as well)
 - They require an **extra i-node** for the link file
- Advantages of symbolic links:
 - There are **no problems with deleting** the original file \Rightarrow the file simply does not exist any more
 - They can **cross the boundaries of machines**, i.e. the linked file can be located on a different machine

<https://powcoder.com>
Add WeChat powcoder

l-nodes

Hard and Soft Links in Unix

```
[pszit@luthien ~]$ pwd
```

```
/home/pszit
```

```
[pszit@luthien ~]$ ls -i labs/
```

```
3250013 req1b.c
```

```
[pszit@luthien ~]$ ln labs/req1b.c hardLink
```

```
[pszit@luthien ~]$ ln -s labs/req1b.c softLink
```

```
[pszit@luthien ~]$ ls -ali hardLink softLink
```

```
3250013 hardLink
```

```
3250021 softLink -> labs/req1b.c
```

```
[pszit@luthien ~]$ rm labs/req1b.c
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Log Structured File System

Context

Assignment Project Exam Help

- Consider the **creation of a new file** on a Unix system:
 - 1 Allocate, initialise and write the i-node for the file
 - i-nodes are usually located at the start of the disk
 - 2 Update and write the directory entry for the file (directories are tables/files that map names onto i-nodes in Unix)
 - 3 Write the data to the disk
- The corresponding blocks are not necessarily in **adjacent locations!**
- Also in linked lists/FAT file systems blocks can be **distributed all over the disk**

Log Structured File System

Context

Assignment Project Exam Help

- Due to seek and rotational delays, **hard disks are slow** compared to other components in a computer (e.g. CPU, main memory)
 - Can we develop a file system that copes better with the inherent delays of traditional disks?
- A **log structured file system** aims to improve speed of a file system on a traditional hard disk by **minimising head movements** and **rotational delays** using the entire disk as a great big log
- A log is a data structure that is written only at the end

<https://powcoder.com>

Add WeChat powcoder

Log Structured File System

Concept

Assignment Project Exam Help

- Log structured file systems **buffer read and write operations** (i-nodes, data, etc.) in memory, enabling us to **write “larger volumes”** in one go
- Once the buffer is full it is “flushed” to the disk and written as **one contiguous segment** at the end of “a log”
 - I-nodes and data are all written to the same segment
 - Finding i-nodes (traditionally located at the start of the partition) becomes more difficult
- An **i-node map** is maintained in memory to quickly find the address of i-nodes on the disk


<https://powcoder.com>

Add WeChat powcoder

Log Structured File System

Structure

Traditional Unix File System



Log-Structured File System




Figure: Blocks written to create two 1-block files: dir1/file1 and dir2/file2

Log Structured File System

Concept

- How do we deal with deleted files?
- A **cleaner thread** is running in the background and spends its time scanning the log circularly and compacting it.
- A hard drive is treated as a **circular buffer**
- It removes deleted files and files being used right now are marked as free segments as they will be later written at the end.



Figure: Log Structured File System

Log Structured File System

Advantages and Disadvantages

Assignment Project Exam Help

- It greatly **increases disk performance** on writes, file creates, deletes
- Writes are **more robust** as they are done as a single operation (Multiple small writes are more likely to expose the file system to serious inconsistency).
- However, it has not been widely used because it **is highly incompatible** with existing file systems.
- In addition, the cleaner thread takes additional CPU time

<https://powcoder.com>
Add WeChat powcoder

File System Implementations

Journaling File Systems: Example

Assignment Project Exam Help

- **Deleting a file** consists of the following actions:

- 1 Remove the file's directory entry
- 2 Add the file's i-node to the pool of free i-nodes
- 3 Add the file's disk blocks to the free list

- Where can it **go wrong**?, e.g.:

- Directory entry has been deleted and a crash occurs \Rightarrow i-nodes and disk blocks become inaccessible
- The directory entry and i-nodes have been released and a crash occurs \Rightarrow disk blocks become inaccessible

File System Implementations

Journaling File Systems: Example

Assignment Project Exam Help

- **Changing the order** of the events does not necessarily resolve the issues
- Journaling file systems aim at **increasing the resilience** of file systems against **crashes** by recording each update to the file system as a transaction

<https://powcoder.com>

Add WeChat powcoder

File System Implementations

Journaling File Systems: Concept

Assignment Project Exam Help

- The key idea behind a journaling file system is to **log all events** (transactions) before they take place
 - Write the actions that should be undertaken to a log file
 - Carry them out
 - Remove/commit the entries once completed
- If a crash happens in the middle of an action (e.g., deleting a file) the **entry in the log file will remain present after the crash**
- The log can be **examined after the crash** and used to restore the consistency of the file system
- **NTFS** and **EXT3-4** are examples of journaling file systems

<https://powcoder.com>

Add WeChat powcoder

File System Implementations

Virtual File Systems: Concept

Assignment Project Exam Help

- **Multiple file systems** usually coexist on the same computer (e.g., NTFS and ISO9660 for a CD-ROM, NFS)
- These file systems can be **seamlessly integrated** by the operating system (e.g. Unix / Linux)
- This is usually achieved by using **virtual file systems** (VFS)
- VFS relies on **standard object oriented principles** (or manual implementations thereof), e.g. polymorphism

<https://powcoder.com>

Add WeChat powcoder

File System Implementations

Virtual File Systems: Concept

Assignment Project Exam Help

- Consider some code that you are writing, **reading “data records”** (DataObject) from a file
- These records can be stored in **CSV file, or XML File**
- How would you make your code resilient against **changes in the underlying data structure?**

Add WeChat powcoder

File System Implementations

Virtual File Systems: Concept

Assignment Project Exam Help

- Consider some code that you are writing, **reading “data records”** (DataObject) from a file
- These records can be stored in **CSV file, or XML File**
- How would you make your code resilient against **changes in the underlying data structure?**
 - You would hide the **implementation** behind **interfaces** using **Data Access Objects (DAOs)**

File System Implementations

Virtual File Systems: Concept

Assignment Project Exam Help

- We can define a **generic interface**, e.g. `DataReader`, containing a method `public DataObject readData();`
 - In the case of file systems, this would be the **POSIX interface** containing reads, writes, close etc.
- You would hide the CSV and XML code in **specific implementations** of the `DataReader` interface, e.g. `CSVDataReader` and `XMLDataReader`
 - In the case of file systems this would be the file system implementations

<https://powcoder.com>

Add WeChat powcoder

File System Implementations

Virtual File Systems: Concept (Cont'd)

Assignment Project Exam Help

- You would rely on **polymorphism** to call the correct method

- `DataReader dr = new CSVDataReader();`
`dr.readData() // reads data from CSV`
- `DataReader dr = new XMLDataReader();`
`dr.readData() // reads data from XML`

Add WeChat powcoder

File System Implementations

Virtual File Systems: Concept

```
1 import java.io.BufferedReader;
2 import java.io.IOException;
3 import java.io.InputStreamReader;
4
5 public class ClientApplication {
6     public static void main(String[] args) throws IOException {
7         System.out.println("Choose CSV or XML?");
8         BufferedReader br
9             = new BufferedReader(new InputStreamReader(System.in));
10        String type = br.readLine();
11        br.close();
12
13        DataReader reader = null;
14        if (type.equals("CSV")) {
15            reader = new CSVDataReader();
16        } else if (type.equals("XML")) {
17            reader = new XMLDataReader();
18        }
19
20        if (reader != null)
21            System.out.println(reader.readData());
22    }
23 }
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

File System Implementations

Virtual File Systems: Concept

Assignment Project Exam Help

- In a similar way, Unix and Linux **unify different file systems** and present them as a single hierarchy and hides away / abstracts the implementation specific details for the user
- The VFS presents a unified interface to the “outside”
- File system specific code is dealt with in an **implementation layer** that is clearly separated from the interface

<https://powcoder.com>

Add WeChat powcoder

File System Implementations

Virtual File Systems: Concept

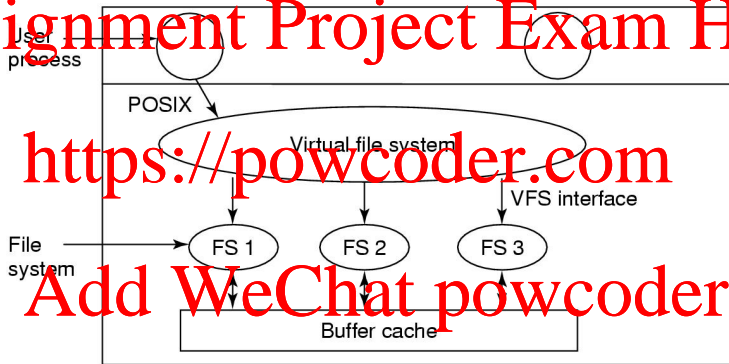


Figure: Virtual File System (Tanenbaum)

File System Implementations

Virtual File Systems: Concept

Assignment Project Exam Help

- The VFS interface commonly contains the **POSIX system calls** (`open`, `close`, `read`, `write`, ...)
- Each file system that meets the VFS requirements **provides an implementation** for the system calls contained in the interface
- Note that implementations can be for **remote file systems** (e.g. `sshfs`), i.e. the file can be stored on a different machine

<https://powcoder.com>
Add WeChat powcoder

File System Implementations

Virtual File Systems: In practice

Assignment Project Exam Help

- Every file system, including the root file system, is **registered with the VFS**
 - A list / table of addresses to the **VFS function calls** (i.e. function pointers) for the specific file system is provided
 - **Every VFS function call** corresponds to a specific **entry in the VFS function table** for the given file system
 - The **VFS maps / translates** the POSIX call onto the “native file system call”
- A virtual file system is essentially **good programming practice**

<https://powcoder.com>

Add WeChat powcoder

File System Implementations

Virtual File Systems: Example

Assignment Project Exam Help

```
$ df -Th
```

Filesystem	Type	Size	Used	Avail	Use%	Mounted on
devtmpfs	devtmpfs	32G	0	32G	0%	/dev
/dev/sda5	btrfs	215G	82G	132G	39%	/
tmpfs	tmpfs	32G	8.6M	32G	1%	/tmp
/dev/sda3	btrfs	215G	82G	132G	39%	/home
/dev/sda1	ext4	477M	166M	282M	38%	/boot
/dev/sdb1	ext4	917G	165G	709G	19%	/data
pszit@bann:	fuse.sshfs	700M	282M	418M	41%	/mnt/bann

Summary

Take-Home Message¹

Assignment Project Exam Help

- Hard and Soft Links
- File system paradigms:
 - Logs: store everything as close as possible
 - Journaling: apply the transaction principle
 - VFS: apply good software design (polymorphism)

<https://powcoder.com>
Add WeChat powcoder

¹Tanenbaum Chapter 4, Sections 4.3.5, 4.3.6, 4.3.7

File systems

Food for thought

Assignment Project Exam Help

Following up the previous question on FAT-32, we mentioned that FAT-32 has severe limitations (e.g. the file allocation table could be too big).

- Why do you think that this file system is still in use in most of flash drives, cameras or MP3 players?
- When would you consider formatting your flash drive with NTFS or ext3?
- If you format your flash drive with (Windows) NTFS, will it work (directly) in Unix systems?

<https://powcoder.com>

Add WeChat powcoder

Submit your answer at:

<https://b.socrative.com/login/student/>

Room name: G52OSC