

Assignment Project Exam Help

Operating Systems and Concurrency

Lecture 12: Memory Management I
G52OSC

<https://powcoder.com>

Add WeChat powcoder
Geert De Maere and Isaac Triguero
{Geert.DeMaere,Isaac.Triguero}@nottingham.ac.uk

University Of Nottingham
United Kingdom

2018

Assignment Project Exam Help

- Contact details:

- Name: Isaac Triguero
- Office: C82 (office hour 10:00-11:00 on Tuesdays - please confirm by e-mail beforehand)
- E-mail: Isaac.Triguero@nottingham.ac.uk

- About me:

- Graduated in 2009, Msci in Computer Science and Engineering
- Completed my PhD in CS in 2014 (Data mining)
- Member of ASAP
- 4 years experience as SysAdmin.
- Specific interest in Data Science

<https://powcoder.com>
Add WeChat powcoder

Remember

Subjects We Will Discuss

Subject	# Lectures	By
Introduction to operating systems/computer design	1-2	GDM
Processes, process scheduling, threading, ...	3-4	GDM
Concurrency (deadlocks)	4-5	GDM
Revision	1	GDM
Memory management, swapping, virtual memory, ...	5-6	IT
File Systems, file structures, management, ...	4-5	IT
Virtualisation and the Cloud	1-2	IT
Revision	1	IT

Table: Module Structure

Important: From now on, we don't have more lectures on **Thursdays at 4pm!**

Assignment Project Exam Help

- 1 Introduction to memory management
- 2 **Modelling** of multi-programming
- 3 Memory management based on **fixed partitioning**

<https://powcoder.com>

Add WeChat powcoder

Memory Management

Memory Hierarchies

Assignment Project Exam Help

- Computers typically have memory hierarchies:
 - **Registers**, L1/L2/L3 **cache**
 - **Main memory**
 - **Disks**
- “**Higher memory**” is faster, more expensive and volatile, “**lower memory**” is slower, cheaper, and non-volatile
- The operating system provides a **memory abstraction**
- Memory can be seen as one **linear array** of bytes/words

<https://powcoder.com>

Add WeChat powcoder

Memory Management

OS Responsibilities

Assignment Project Exam Help

- **Allocate/deallocate** memory when requested by processes, keep track of **used/unused** memory
- **Distribute memory** between processes and simulate an “**infinitely large**” memory space
- **Control access** when multiprogramming is applied
- **Transparently** move data from **memory** to **disk** and vice versa

<https://powcoder.com>

Add WeChat powcoder

Memory Management

History of Memory Management

Assignment Project Exam Help

- Memory management has **evolved**: memory management in **modern computers** is very different from **early computers**
- **History repeats itself**:
 - Many of the **early ideas underpin** more **modern memory management** approaches (e.g. relocation)
 - **Modern consumer electronics** often require **less complex memory management** approaches

<https://powcoder.com>

Add WeChat powcoder

Partitioning

Approaches: Contiguous vs. Non-Contiguous



Figure: Contiguous




Figure: Non-contiguous

Models

Approaches: Contiguous vs. Non-Contiguous

Assignment Project Exam Help

- **Contiguous memory management models** allocate memory in **one single block** without any holes or gaps
- **Non-contiguous memory management models** are capable of allocating memory in **multiple blocks**, or **segments**, which may be **placed anywhere in physical memory** (i.e., not necessarily next to each other)

<https://powcoder.com>
Add WeChat powcoder

Partitioning

Contiguous Approaches

Assignment Project Exam Help

- **Mono-programming**: one single partition for user processes
- **Multi-programming with fixed partitions**
 - Fixed **equal** sized partitions
 - Fixed **non-equal** sized partitions
- **Multi-programming with dynamic partitions**

<https://powcoder.com>
Add WeChat powcoder

Mono-Programming

No Memory Abstraction

- Only one single user process is in memory/executed at any point in time (no multi-programming)
- A fixed region of memory is allocated to the OS/kernel, the remaining memory is reserved for a **single process** (MS-DOS worked this way)
- This process has **direct access** to **physical memory** (i.e. no address translation takes place)




Figure: Mono-programming

Mono-Programming

No Memory Abstraction: Properties

Assignment Project Exam Help

- Every process is allocated **contiguous block of memory**, i.e. it contains no “holes” or “gaps” (\Leftrightarrow **non-contiguous allocation**)
- **One process** is allocated the **entire memory space**, and the process is **always located in the same address space**
- **No protection** between different user processes required (one process)
- **Overlays**¹ enable the programmer to use more memory than available (burden on programmer)

¹Tanenbaum - Section 3.3 (page 194)

Mono-Programming

No Memory Abstraction: Properties (Cont'ed)

Assignment Project Exam Help

- **Shortcomings of mono-programming:**

- Since a process has **direct access to the physical memory**, it may have **access to OS memory**
- The operating system can be seen as a process - so we have **two processes anyway**
- **Low utilisation** of hardware resources (CPU, I/O devices, etc.)
- Mono-programming is unacceptable as **multiprogramming is expected** on modern machines

- Direct memory access and mono-programming are common in basic **embedded systems** and **modern consumer electronics**, e.g. washing machines, microwaves, car's ECUs, etc.

Mono-Programming

Simulating Multi-Programming

Assignment Project Exam Help

- Simulate multi-programming through **swapping**
 - **Swap process** out to the disk and load a new one (context switches would become **time consuming**)
 - Apply **threads** within the same process (limited to one process)
- Assumption that **multiprogramming** can **improve CPU utilisation?**
 - Intuitively, this is true
 - How do we model this?

<https://powcoder.com>

Add WeChat powcoder

Multi-Programming

A Probabilistic Model

- There are n processes in memory
- A process spends p percent of its time waiting for I/O
- **CPU Utilisation** is calculated as 1 minus the time that all processes are waiting for I/O: e.g., $p = 0.9$ then CPU utilisation = $1 - 0.9 \Rightarrow 0.1$ ($1 - p$)
- The probability that **all n processes are waiting for I/O** (i.e., the CPU is idle) is p^n , i.e. $p \times p \times p \dots$
- The **CPU utilisation** is given by $1 - p^n$



Multi-Programming

A Probabilistic Model

Assignment Project Exam Help

- With an **I/O wait time of 20%**, almost **100% CPU utilisation** can be achieved with four processes ($1 - 0.2^4$)
- With an **I/O wait time of 90%**, 10 processes can achieve about **65% CPU utilisation** ($1 - 0.9^{10}$)
- **CPU utilisation goes up with the number of processes and down for increasing levels of I/O**

<https://powcoder.com>

Add WeChat powcoder

Multi-Programming

A Probabilistic Model

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Processes	I/O Ratio		
	0.2	0.5	0.8
1	0.80	0.50	0.20
2	0.96	0.75	0.36
3	0.99	0.88	0.49
4	1.00	0.94	0.59
5	1.00	0.97	0.67
6	1.00	0.98	0.74
7	1.00	0.99	0.79
8	1.00	1.00	0.83
9	1.00	1.00	0.87
10	1.00	1.00	0.89

Table: CPU utilisation as a function of the I/O ratio and the number of processes


Multi-Programming

A Probabilistic Model

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Multi-Programming

A Probabilistic Model

- Assume that:
 - A computer has **one megabyte of memory**
 - The **OS takes up 200k**, leaving room for **four 200k processes**
- Then:
 - If we have an I/O wait time of **80%**, then we will achieve just under **60% CPU utilisation** ($1 - 0.8^4$)
 - If we add **another megabyte of memory**, it would allow us to run another **five processes**
 - We can achieve about **87% CPU utilisation** ($1 - 0.8^9$)
 - If we add another **megabyte of memory** (fourteen processes) we will find that the CPU utilisation will increase to **about 96%** ($1 - 0.8^{14}$)

Multi-Programming

A Probabilistic Model

- Assume that:
 - A computer has **one megabyte of memory**
 - The **OS takes up 200k**, leaving room for **four 200k processes**
- Then:
 - If we have an I/O wait time of **80%**, then we will achieve just under **60% CPU utilisation** ($1 - 0.8^4$)
 - If we add **another megabyte of memory**, it would allow us to run another **five processes**
 - We can achieve about **87% CPU utilisation** ($1 - 0.8^9$)
 - If we add another **megabyte of memory** (fourteen processes) we will find that the CPU utilisation will increase to **about 96%** ($1 - 0.8^{14}$)
- **Multi-programming** does enable to **improve resource utilisation**
 ⇒ memory management should provide support for multi-programming

Multi-Programming

A Probabilistic Model

Assignment Project Exam Help

Caveats:

- This model assumes that **all processes are independent**, this is not true
- More complex models could be built using **queueing theory**, but we can still use this simplistic model to **make approximate predictions**

Add WeChat powcoder


<https://powcoder.com>

Partitioning

Fixed Partitions of equal size

Assignment Project Exam Help

- Divide memory into **static, contiguous** and **equal sized partitions** that have a **fixed size and fixed location**
 - Any process can take **any** (large enough) **partition**
 - Allocation of **fixed equal sized partitions** to processes is **trivial**
 - **Very little overhead and simple implementation**
 - The OS keeps track of which partitions are being **used** and which are **free**



(a) Equal-size partitions

Figure: From Stallings


Partitioning

Fixed Partitions of equal size

Assignment Project Exam Help

- **Disadvantages of static equal-sized partitions**

- **Low memory utilisation and internal fragmentation:** partition may be unnecessarily large
- **Overlays must be used** if a program does not fit into a partition (burden on programmer)



(a) Equal-size partitions

Figure: From Stallings


Partitioning

Fixed Partitions of non-equal size

Assignment Project Exam Help

- Divide memory into **static and non-equal sized partitions** that have a **fixed size** and **fixed location**

- **Reduces internal fragmentation**
- The **allocation of processes to partitions** must be **carefully considered**



(b) Unequal-size partitions

Figure: From Stallings

Partitioning

Fixed Partitions (Allocation Methods)

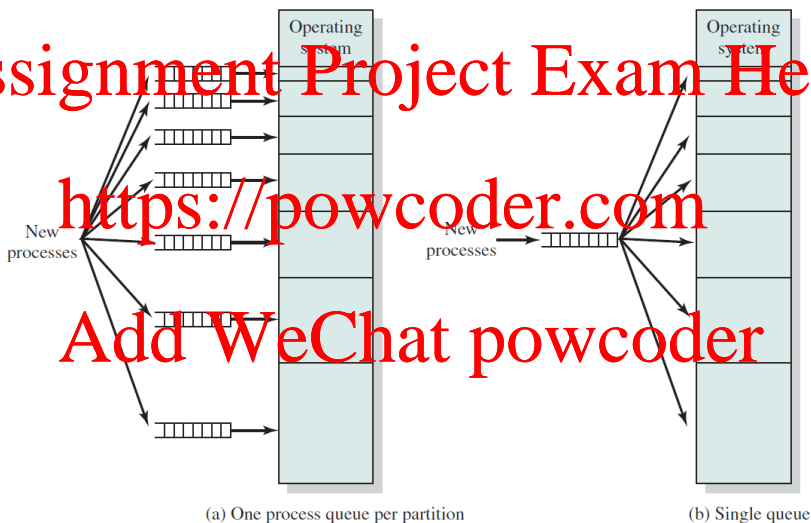


Figure: From Stallings

Partitioning

Fixed Partitions (Allocation Methods)

Assignment Project Exam Help

- One **private queue** per partition:
 - Assigns each process to the **smallest partition** that it would fit in
 - Reduces **internal fragmentation**
 - Can **reduce memory utilisation** (e.g., lots of small jobs result in unused large partitions) and result in **starvation**
- A single **shared queue** for all partitions can allocate small processes to **large partitions** but results in **increased internal fragmentation**

Assignment Project Exam Help

- **Mono-programming** and **absolute addressing** (no memory abstraction)
- Why multi-programming: CPU utilisation modelling.
- Memory management for **Multi-programming**: fixed (non-)equal partitions.

<https://powcoder.com>
Add WeChat powcoder

²Resources: Tanenbaum Section 3.1, 3.2. Stallings Section 7.1, 7.2