

Assignment Project Exam Help

Operating Systems and Concurrency

Lecture 15: Memory Management IV
G52OSC

<https://powcoder.com>

Add WeChat powcoder
Geert De Maere and Isaac Triguero
{Geert.DeMaere,Isaac.Triguero}@Nottingham.ac.uk

University Of Nottingham
United Kingdom

2018

Assignment Project Exam Help

- The **principles of paging** are:
 - Main memory is divided into small equal sized **frames**
 - Each process is divided into **pages** of equal size
 - A page table contains multiple "relocation registers" (**page table**) to map the pages on to frames
- The **benefits** of paging include:
 - Reduced internal fragmentation
 - No **external fragmentation**

Assignment Project Exam Help

- Address Translation/Implementation (revisited)
- Principles behind **virtual memory**
- Complex/large **page tables**

<https://powcoder.com>

Add WeChat powcoder

Recall

Memory as a linear array

- Memory can be seen as one **linear array** of **bytes**/words: *byte memory* [N];

- Address ranges from 0 through (N - 1)

- N address lines can be used to specify 2^N distinct addresses

- E.g. in a 16-bit machine, we can address up to 2^{16} distinct addresses. Nowadays, **each cell** of memory is typically a **byte**. Thus, we can address up to 64 kilobytes.

- If the memory is split into 16 blocks ($=2^4$);
Block size = $2^{16}/2^4 = 2^{12} = 4$ kilobytes




Figure: Memory as a linear array of bytes (16-bit)

Paging

Relocation: Address Translation

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder




Figure: Address Translation

Paging

Address Translation: Implementation

- A **logical (physical) address** is relative to the start of the **program (memory)** and consists of two parts:
 - The **right most n bits** that represent the **offset within the page frame**
 - e.g. 12 bits for the offset, allowing up to 4096 (2^{12}) bytes per page (frame)
 - The **left most n bits** that represent the **page (frame) number**
 - e.g. 4 bits for the page number allowing 16 (2^4) pages (frames)
- The **offset within the page and frame remains the same** (they are the same size)
- The page number to frame number mapping is held in the **page table**




Figure: Logical Address

Paging

Address Translation: Implementation

- A **logical (physical) address** is relative to the start of the **program (memory)** and consists of two parts:
 - The **right most n bits** that represent the **offset within the page frame**
 - e.g. 12 bits for the offset, allowing up to 4096 (2^{12}) bytes per page (frame)
 - The **left most n bits** that represent the **page (frame) number**
 - e.g. 4 bits for the page number allowing 16 (2^4) pages (frames)
- The **offset within the page and frame remains the same** (they are the same size)
- The page number to frame number mapping is held in the **page table**

Add WeChat powcoder



Figure: Logical Address

Paging

Address Translation: Implementation

- A **logical (physical) address** is relative to the start of the **program (memory)** and consists of two parts:
 - The **right most n bits** that represent the **offset within the page frame**
 - e.g. 12 bits for the offset, allowing up to 4096 (2^{12}) bytes per page (frame)
 - The **left most n bits** that represent the **page (frame) number**
 - e.g. 4 bits for the page number allowing 16 (2^4) pages (frames)
- The **offset within the page and frame remains the same** (they are the same size)
- The page number to frame number mapping is held in the **page table**

Add WeChat powcoder




Figure: Logical Address

Paging

Address Translation: Implementation

- A **logical (physical) address** is relative to the start of the **program (memory)** and consists of two parts:
 - The **right most n bits** that represent the **offset within the page frame**
 - e.g. 12 bits for the offset, allowing up to 4096 (2^{12}) bytes per page (frame)
 - The **left most n bits** that represent the **page (frame) number**
 - e.g. 4 bits for the page number allowing 16 (2^4) pages (frames)
- The **offset within the page and frame remains the same** (they are the same size)
- The page number to frame number mapping is held in the **page table**

Add WeChat powcoder




Figure: Logical Address

Paging

Address Translation: Implementation

- A **logical (physical) address** is relative to the start of the **program (memory)** and consists of two parts:
 - The **right most n bits** that represent the **offset within the page frame**
 - e.g. 12 bits for the offset, allowing up to 4096 (2^{12}) bytes per page (frame)
 - The **left most n bits** that represent the **page (frame) number**
 - e.g. 4 bits for the page number allowing 16 (2^4) pages (frames)
- The **offset within the page and frame remains the same** (they are the same size)
- The page number to frame number mapping is held in the **page table**



Figure: Logical Address

Paging

Address Translation: Implementation

- A **logical (physical) address** is relative to the start of the **program (memory)** and consists of two parts:
 - The **right most n bits** that represent the **offset within the page frame**
 - e.g. 12 bits for the offset, allowing up to 4096 (2^{12}) bytes per page (frame)
 - The **left most n bits** that represent the **page (frame) number**
 - e.g. 4 bits for the page number allowing 16 (2^4) pages (frames)
- The **offset within the page and frame remains the same** (they are the same size)
- The page number to frame number mapping is held in the **page table**

Add WeChat powcoder



Figure: Logical Address

Paging

Address Translation: Implementation

- A **logical (physical) address** is relative to the start of the **program (memory)** and consists of two parts:
 - The **right most n bits** that represent the **offset within the page frame**
 - e.g. 12 bits for the offset, allowing up to 4096 (2^{12}) bytes per page (frame)
 - The **left most n bits** that represent the **page (frame) number**
 - e.g. 4 bits for the page number allowing 16 (2^4) pages (frames)
- The **offset within the page and frame remains the same** (they are the same size)
- The page number to frame number mapping is held in the **page table**



Figure: Logical Address

Paging

Address Translation: Implementation

- A **logical (physical) address** is relative to the start of the **program (memory)** and consists of two parts:
 - The **right most n bits** that represent the **offset within the page frame**
 - e.g. 12 bits for the offset, allowing up to 4096 (2^{12}) bytes per page (frame)
 - The **left most n bits** that represent the **page (frame) number**
 - e.g. 4 bits for the page number allowing 16 (2^4) pages (frames)
- The **offset within the page and frame remains the same** (they are the same size)
- The page number to frame number mapping is held in the **page table**

Add WeChat powcoder



Figure: Logical Address

Paging

Address Translation: Implementation

- A **logical (physical) address** is relative to the start of the **program (memory)** and consists of two parts:
 - The **right most n bits** that represent the **offset within the page frame**
 - e.g. 12 bits for the offset, allowing up to 4096 (2^{12}) bytes per page (frame)
 - The **left most n bits** that represent the **page (frame) number**
 - e.g. 4 bits for the page number allowing 16 (2^4) pages (frames)
- The **offset within the page and frame remains the same** (they are the same size)
- The page number to frame number mapping is held in the **page table**

Add WeChat powcoder



Figure: Logical Address

Paging

Address Translation: Implementation

- A **logical (physical) address** is relative to the start of the **program (memory)** and consists of two parts:
 - The **right most n bits** that represent the **offset within the page frame**
 - e.g. 12 bits for the offset, allowing up to 4096 (2^{12}) bytes per page (frame)
 - The **left most n bits** that represent the **page (frame) number**
 - e.g. 4 bits for the page number allowing 16 (2^4) pages (frames)
- The **offset within the page and frame remains the same** (they are the same size)
- The page number to frame number mapping is held in the **page table**

Add WeChat powcoder



Figure: Logical Address

Paging

Address Translation: Implementation

- A **logical (physical) address** is relative to the start of the **program (memory)** and consists of two parts:
 - The **right most n bits** that represent the **offset within the page frame**
 - e.g. 12 bits for the offset, allowing up to 4096 (2^{12}) bytes per page (frame)
 - The **left most n bits** that represent the **page (frame) number**
 - e.g. 4 bits for the page number allowing 16 (2^4) pages (frames)
- The **offset within the page and frame remains the same** (they are the same size)
- The page number to frame number mapping is held in the **page table**



Figure: Logical Address

Paging

Address Translation: Implementation

- A **logical (physical) address** is relative to the start of the **program (memory)** and consists of two parts:
 - The **right most n bits** that represent the **offset within the page frame**
 - e.g. 12 bits for the offset, allowing up to 4096 (2^{12}) bytes per page (frame)
 - The **left most n bits** that represent the **page (frame) number**
 - e.g. 4 bits for the page number allowing 16 (2^4) pages (frames)
- The **offset within the page and frame remains the same** (they are the same size)
- The page number to frame number mapping is held in the **page table**



Figure: Logical Address

Paging

Address Translation: Implementation

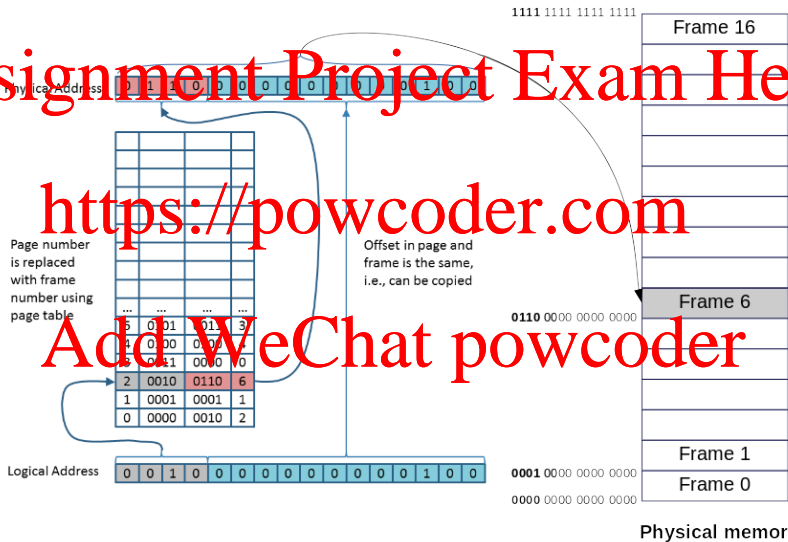


Figure: Address Translation

Assignment Project Exam Help

- Steps in **address translation**:

- ① **Extract the page number** from logical address
- ② Use page number as an index to **retrieve the frame number** in the page table
- ③ Add the “**logical offset within the page**” to the start of the physical frame

- **Hardware implementation** of address translation

- ① The CPU's **memory management unit (MMU)** intercepts logical addresses
- ② MMU uses a page table as above
- ③ The resulting **physical address** is put on the **memory bus**

Assignment Project Exam Help

- Are there any other **benefits of paging**?
 - Code execution and data manipulation are usually **restricted to a small subset** (i.e. limited number of pages) at any point in time
 - **i.e. code and data references within a process are usually clustered**
⇒ This is called the **principle of locality**
- **Not all pages** have to be **loaded** in memory at the same time ⇒ **virtual memory**
 - Loading an entire set of pages for an entire program/data set into memory is **wasteful**
 - Desired blocks could be **loaded on demand**

Virtual Memory

An Example




Figure: Virtual Memory

Assignment Project Exam Help

- The **resident set** refers to the pages that are loaded in main memory
- A **page fault** is generated if the processor accesses a page that is **not in memory**
 - A page fault results in an **interrupt** (process enters **blocked state**)
 - An **I/O operation** is started to bring the missing page into main memory
 - A **context switch** (may) take place
 - An **interrupt signal** that the I/O operation is complete (process enters **ready state**)

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

1. Trap operating system

- Save registers/process state
- Analyse interrupt (i.e., identify page fault)
- Validate page reference, determine page location
- Issue disk I/O: queueing, seek, latency, transfer

2. Context switch (optional)

3. Interrupt for I/O completion

- Store process state/registers
- Analyse interrupt from disk
- Update page table (page in memory)
- Wait for original process to be scheduled

4. Context switch to original process

<https://powcoder.com>

Add WeChat powcoder

Virtual Memory

The Benefits

- Being able to maintain **more processes** in main memory through the use of virtual memory **improves CPU utilisation**
 - Individual processes take up less memory since they are only partially loaded
- Virtual memory allows the **logical address space** (i.e processes) to be **larger than physical address space** (i.e. main memory)
 - 64 bit machine $\Rightarrow 2^{64}$ logical addresses (theoretically)

Add WeChat powcoder



Figure: Logical Address \Rightarrow physical address

Virtual Memory

The Benefits

- Being able to maintain **more processes** in main memory through the use of virtual memory **improves CPU utilisation**
 - Individual processes take up less memory since they are only partially loaded
- Virtual memory allows the **logical address space** (i.e processes) to be **larger than physical address space** (i.e main memory)
 - 64 bit machine $\Rightarrow 2^{64}$ logical addresses (theoretically)

Add WeChat powcoder



Figure: Logical Address \Rightarrow physical address

Virtual Memory

Page Tables Revisited: Contents of a Page Entry

- A “**present/absent bit**” that is set if the page/frame is in memory
- A “**modified bit**” that is set if the page/frame has been modified (only modified pages have to be written back to disk when evicted)
- A “**referenced bit**” that is set if the page is in use
- **Protection and sharing bits** read, write, execute or combinations thereof

<https://powcoder.com>

Add WeChat powcoder

...	0/1	0/1	RWX	0/1	Frame Number
-----	-----	-----	-----	-----	--------------

Figure: Page Table Entry

Assignment Project Exam Help

- On a **16 bit machine**, the total address space is 2^{16}

- Assuming that 10 bits are used for the offset (2^{10})

- 6 bits can be used to number the pages

- I.e. $2^6 = 64$ pages can be maintained

- In a **32 bit machine**, total address space is 2^{32}

- Assuming pages of 2^{12} bits (4KB)

- 20 bits can be used to number the pages

- I.e. 2^{20} pages (approx. 1 million) can be maintained

- On a **64 bit machine** ...

<https://powcoder.com>

Add WeChat powcoder

Virtual Memory

Page Tables Revisited: Dealing with Large Page Tables

Assignment Project Exam Help

- How do we deal with **the increasing size of page tables**, i.e., where do we store them?
 - Their size prevents them from being **stored in registers**
 - They have to be stored in (virtual) **main memory**:
 - **Multi-level** page tables
 - **Inverted page tables** (for large virtual address spaces)
- How can we maintain **acceptable speeds**? : address translation happens at every memory reference, it has to be fast!
 - Accessing main memory results in **memory stalls**

Virtual Memory

Page Tables Revisited: Multi-level Page Tables

Assignment Project Exam Help

- **Solution:** Page the page table!
- We keep tree-like structures to hold page tables
- Divide the page number into
 - An index to a page table of second level
 - A page within a second level page table
- No need to keep all page tables in memory all time




Figure: Multi-level page tables (from Tanenbaum)

Virtual Memory

Page Tables Revisited: Multi-level Page Tables

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

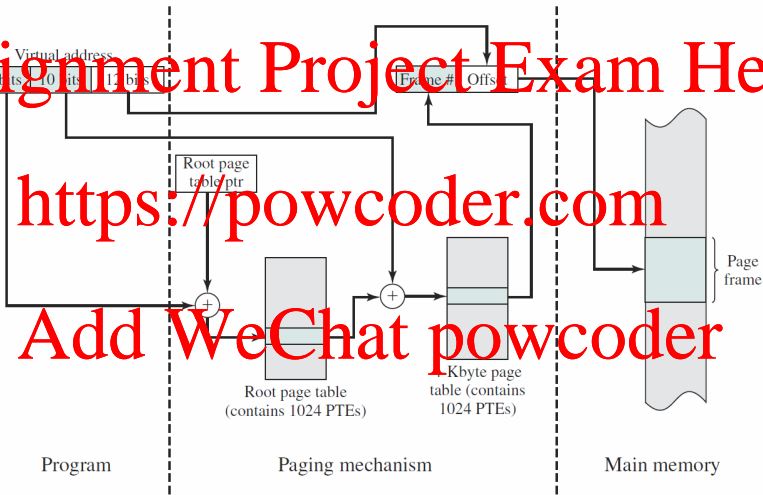


Figure: Multi-level Address Translation (from Stallings)

Assignment Project Exam Help

- **Memory organisation** of multi-level page tables:
 - The **root page table** is always maintained in memory
 - Page tables themselves are maintained in **virtual memory** due to their size
- Assume that a **fetch** from main memory takes T nano seconds
 - With a **single page table level**, access is $2 \times T$
 - With **two page table levels**, access is $3 \times T$
 - ...

<https://powcoder.com>
Add WeChat powcoder

Paging

Address Translation: Exercises

- Given a 4KB page/frame size, and a 16-bit address space, calculate:

- Number **M** of bits for offset within a page.
- Number **N** of bits for representing pages. So, number of pages

- What is the physical address for 0, 8192, 20500 using this page table?

Pages		Frames	
0	0000	0010	2
1	0001	0001	1
2	0010	0110	6
3	0011	0000	0
4	0100	0100	4
5	0101	0011	3
6	0110	X	X
7	0111	X	X
8	1000	X	X
9	1001	0111	5
10	1010	X	X
11	1011	0111	7
12	1100	X	X

Table: Page Table

<https://b.socrative.com/login/student/>

Room name: G52OSC

Assignment Project Exam Help

- **Paging** splits logical and physical address spaces into small **pages/frames** to reduce internal and external fragmentation
- **Virtual memory** exploits the principle of **locality** and allows for processes to be **loaded only partially** into memory, **large logical address spaces** require “different” approaches

<https://powcoder.com>
Add WeChat powcoder

¹Tanenbaum Section 3.3