

Assignment Project Exam Help

Operating Systems and Concurrency

Lecture 6: Process Scheduling

G52CSC/COMP2007

<https://powcoder.com>

Geert De Maere

(Isaac Triguero)

Add WeChat powcoder

{Geert.DeMaere,Isaac.Triguero}@Nottingham.ac.uk

University Of Nottingham
United Kingdom

2018

Assignment Project Exam Help

- 1 Threads vs. processes
- 2 Thread implementations (user, kernel and hybrid)
- 3 PThreads

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

- **Multi-level feedback queues**
- Scheduling in **Windows 7** - illustration
- Scheduling in **Linux** (implementation in labs)
- CPU **affinity** and **load balancing**
- Scheduling related **processes/threads**

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

- Jobs can have **different priority levels** that are **fixed**
- Jobs of the **same priority** are run in **round robin** fashion
- Priority queues are usually implemented by using **multiple queues**, one for each priority level

Add WeChat powcoder

Assignment Project Exam Help

- Different **scheduling algorithms** can be used for the **individual queues** (e.g., round robin, SJF, FCFS)
- **Feedback queues** allow priorities to change dynamically, i.e., jobs can move between queues:
 - Move to **lower priority queue** if too much CPU time is used (prioritise I/O and interactive processes)
 - Move to **higher priority queue** to prevent starvation and avoid inversion of control

Assignment Project Exam Help

Process A (low)

Process B (high)

Process C (high)

• • •

requestX

```

receive X

```

• • •

RUN

• • •

request X

• • •

blocked

• • •

• •

RUN

CC

Assignment Project Exam Help

- Defining characteristics of feedback queues include:
 - The **number of queues**
 - The **scheduling algorithms** used for the individual queues
 - **Migration policy** between queues
 - Initial **access** to the queues
- Feedback queues are highly **configurable** and offer significant flexibility

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

- An **interactive system** using a **preemptive scheduler** with **dynamic priority levels**
 - Two priority classes with 16 different priority levels exist
 - “**Real time**” processes/threads have a **fixed priority level**
 - “**Variable**” processes/threads can have their priorities **boosted temporarily**
- A **round robin algorithm** is used within the queues

<https://powcoder.com>

Add WeChat powcoder

Multi-level Feedback Queues

Windows 7 (Cont'ed)

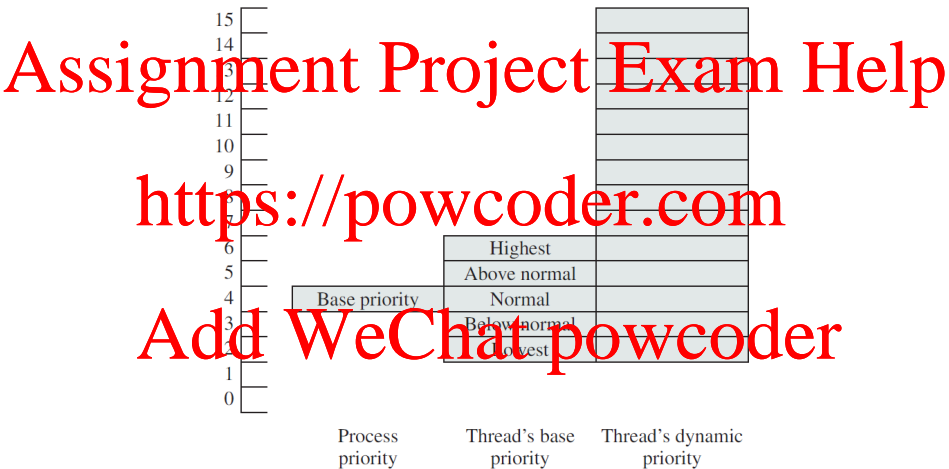


Figure: Priorities in Windows 7 (Stallings, 7th edition)

Assignment Project Exam Help

- Priorities are based on the **process base priority** (between 0-15) and **thread base priority** (± 2 relative to the process priority)
- A thread's **priority dynamically changes** during execution between its base priority and the maximum priority within its class
 - **Interactive I/O bound processes** (e.g. keyboard) receive a **larger boost**
 - Boosting priorities prevents **priority inversion**

<https://powcoder.com>
Add WeChat powcoder

Assignment Project Exam Help

- Code available on Moodle

- Examples:

- 1 Variable process, seven threads, equal priority, 1 core/processor
- 2 Variable process, three threads, different priorities, 1 core/processor
- 3 Real time process, three threads, different priorities, 1 core/processor, admin privileges

<https://powcoder.com>
Add WeChat powcoder

Process scheduling has evolved over different versions of Linux to account for **multiple processors/cores**, processor **affinity**, and **load balancing** between cores

- Linux distinguishes between two types of tasks for scheduling:
 - Real time tasks** (to be POSIX compliant), divided into:
 - Real time FIFO tasks
 - Real time Round Robin tasks
 - Time sharing tasks** using a **preemptive** approach (similar to **variable** in Windows)
- The most recent scheduling algorithm in Linux for **time sharing tasks** is the “**completely fair scheduler**” (CFS, before the 2.6 kernel, this was an $O(1)$ scheduler)

Assignment Project Exam Help

- **Real time FIFO** tasks have the **highest priority** and are scheduled using a **FCFS** approach, using **preemption** if a **higher priority** job shows up
- **Real time round robin** tasks are preemptable by **clock interrupts** and have a **time slice** associated with them
- Both approaches **cannot guarantee hard deadlines**

<https://powcoder.com>
Add WeChat powcoder

Scheduling in Linux

Time Sharing Tasks (equal priority)

Assignment Project Exam Help

- The **SFS** divides the **CPU time** between all processes
- If all N processes have the **same priority**:
 - They will be allocated a “time slice” equal to $\frac{1}{N}$ times the available CPU time
 - I.e., if N equals 5, every process will receive 20% of the processor's time
- The length of the **time slice** and the “available CPU time” are based on the **targeted latency** (\Rightarrow every process should **run at least once** during this interval)
- If N is very large, the **context switch time** will be dominant. Hence a lower bound on the “time slice” is imposed by the minimum granularity
 - A process's time slice can be **no less** than the **minimum granularity** (response time will deteriorate)

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

- A **weighting scheme** is used to take different priorities into account
- If process have **different priorities**:
 - Every process i is allocated a **weight** w_i that reflects its priority
 - The “time slice” allocated to process i is then **proportional to** $\frac{w_i}{\sum_{j \in N} w_j}$
- The tasks with the **lowest proportional amount of “used CPU time”** are selected first

Assignment Project Exam Help

- **Single processor** machine: **which process (thread)** to run next (one dimensional)
- Scheduling decisions on a **multi-processor/core** machine include:
 - Which process (thread) to run **where**, i.e., which CPU?
 - Which process (thread) to run **when**?

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

- A single or multi-level queue **shared** between all CPUs
- Advantage: automatic **load balancing**
- Disadvantages:
 - **Contention** for the queues (locking is needed)
 - *"All CPUs are equal, but some are more equal than others"*: does not account for **processor affinity**:
 - **Cache** becomes invalid when moving to a different CPU
 - Translation look aside buffers (TLBs part of the MMU) become invalid
- Windows will allocate the **highest priority threads** to the individual CPUs/cores

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

- Each CPU has a **private (set) of queues**
- Advantages:
 - **CPU affinity** is automatically satisfied
 - **Contention** for shared queue is minimised
- Disadvantages: less **load balancing**
- **Push and pull migration** between CPUs is possible

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

- **Related:** multiple threads that **communicate** with one another and **ideally run together** (e.g. search algorithm)
- **Unrelated:** e.g. processes threads that are **independent**, possibly started by **different users** running **different programs**

Add WeChat powcoder

Scheduling Related Threads

Working Together

- E.g., threads belong to the same process and are **cooperating**, e.g. they **exchange messages** or **share information**, e.g.

- Process A has thread A_0 and A_1 , A_0 and A_1 cooperate

- Process B has thread B_0 and B_1 , B_0 and B_1 cooperate

- The scheduler selects A_0 and B_1 to run first, then A_1 and B_0 , and A_0 and A_1 , and B_0 and B_1 run on different CPUs

- They try to send messages to the other threads, which are still in the ready state




Figure: Tanenbaum, Chapter 8

Assignment Project Exam Help

- The aim is to get threads **running**, as much as possible, at the **same time across multiple CPUs**
- Approaches include:
 - **Space** sharing
 - **Gang** scheduling

<https://powcoder.com>
Add WeChat powcoder

Assignment Project Exam Help

- Approach:
 - N threads are allocated to N **dedicated CPUs**
 - N threads are kept waiting until N CPUs are available
 - **Non-preemptive**, i.e. blocking calls result in **idle CPUs** (less context switching overhead but results in CPU idle time)
- The number N can be **dynamically adjusted** to match processor capacity

Scheduling Related Threads

Gang scheduling

- Time slices are **synchronised** and the scheduler **groups threads** together to run simultaneously (as much as possible)
- A **preemptive** algorithm (timer interrupts are still used)
- Blocking threads** result in idle CPUs
 - I.e. if a thread blocks, e.g., due to an I/O call, the rest of the time slice will be unused (due to the time slice synchronisation across all CPUs)

	CPU					
	0	1	2	3	4	5
0	A ₀	A ₁	A ₂	A ₃	A ₄	A ₅
1	B ₀	B ₁	B ₂	C ₀	C ₁	C ₂
2	D ₀	D ₁	D ₂	D ₃	D ₄	E ₀
3	E ₁	E ₂	E ₃	E ₄	E ₅	E ₆
4	A ₀	A ₁	A ₂	A ₃	A ₄	A ₅
5	B ₀	B ₁	B ₂	C ₀	C ₁	C ₂
6	D ₀	D ₁	D ₂	D ₃	D ₄	E ₀
7	E ₁	E ₂	E ₃	E ₄	E ₅	E ₆

Assignment Project Exam Help

- Scheduling on **Windows** and **Linux**
- **Multi-processor/core** scheduling is “a bit different” (load balancing, processor affinity, etc.)
 - **Related** and **unrelated** threads
 - **Shared** or **private** queues
 - **Static** scheduling or **dynamic** scheduling

Add WeChat powcoder